

ORACLE

Analyzing 1 billion rows in SQL

And how fast databases can really be

Gerald Venzl

Lead Product Manager of Developer Initiatives

Oracle Database Development



What is this talk about?

We will analyze **1 billion rows** of
synthetical weather station data in
SQL → **LIVE**



Gerald Venzl

- Oracle Database Product Manager
- CNCF Ambassador
 - CNCF.io – *make cloud native computing ubiquitous*
- ISO SQL Standard member
- SQL & Performance enthusiast





Gerald Venzl



github.com/gvenzl



geraldonit.com



linkedin.com/in/geraldvenzl



tobsky.com/gvenzl

The data:

- 1 billion lines in a text file
- 13 GB of data
- The content:

```
$ head measurements_413.txt
```

```
Tamale;34.1
```

```
San Juan;34.0
```

```
St. John's;18.6
```

```
Omaha;3.9
```

```
Lahore;28.4
```

```
Libreville;11.2
```

```
Abéché;38.6
```

```
Naha;31.6
```

```
Phnom Penh;24.3
```

```
Gangtok;21.1
```



What we want

Retrieve temperature measurement values and
calculate the min, average, and max temperature
per weather station in alphabetical order

What that implies

- Scan 13 GB of data
- Group rows per station
- Aggregate min, avg, max values per station
- Sort aggregate

```
SELECT station_name, MIN(val), AVG(val), MAX(val)
FROM <data>
GROUP BY station_name
ORDER BY station_name;
```

In Java

```
public static void main(String[] args) throws IOException {
    Collector<Measurement, MeasurementAggregator, ResultRow> collector = Collector.of(
        MeasurementAggregator::new,
        (a, m) -> {
            a.min = Math.min(a.min, m.value);
            a.max = Math.max(a.max, m.value);
            a.sum += m.value;
            a.count++;
        },
        (agg1, agg2) -> {
            var res = new MeasurementAggregator();
            res.min = Math.min(agg1.min, agg2.min);
            res.max = Math.max(agg1.max, agg2.max);
            res.sum = agg1.sum + agg2.sum;
            res.count = agg1.count + agg2.count;

            return res;
        },
        agg -> {
            return new ResultRow(agg.min, (Math.round(agg.sum * 10.0) / 10.0) / agg.count, agg.max);
        });

    Map<String, ResultRow> measurements = new TreeMap<>(Files.lines(Paths.get(FILE))
        .map(l -> new Measurement(l.split(";")))
        .collect(groupingBy(m -> m.station(), collector)));
}
```



What we will be using

- 128 cores of *Intel Xeon Platinum 8358 CPU @ 2.60GHz (max turbo 3.4GHz)*
 - Oracle Cloud Infrastructure [BM.Standard3.64](#)
- 1 TB of RAM
- 2 TB of disk space

What we will be testing

DuckDB 1.1.3



Postgres 17



Oracle 19c



Let's get started

github.com/gvenzl/one-billion-rows-database

ORACLE

Our mission is to help people see
data in new ways, discover insights,
unlock endless possibilities.

