



# **STOP ALERT FATIGUE MONITOR WHAT MATTERS**


**January 2025**



Quality and Reliability  
Lead at RocketReach

Tester, Breaker, Fixer  
of all things

# REWIND



# 2022 IN REVIEW



**Sev I Incidents  
Weekly**



**Alerts Triggered  
100-1000 per  
day**



**Estimated Downtime  
8-20 hours a  
month**

# FAST FORWARD



# 2024 IN REVIEW



**Sev I Incidents  
Every Other  
Month**



**Alerts Triggered  
10-100 a day**



**Estimated Downtime  
1-4 hours a  
month**

# COMPARING 2022 TO 2024

Weekly to  
Monthly  
Sev I Incidents

---



1000s to Double  
Digits  
Total Alerted Incidents

---



99.4%  
uptime  
Total Downtime

---







# HOW DID WE GET THERE



# HOW DID WE GET THERE



Things happened,  
lets fix them

**REACTIONARY**



# HOW DID WE GET THERE

8

Reply hazy, try again  
later

**PRECAUTIONARY**

# HOW DID WE GET THERE



"The best time to start  
was last year"

**TIME AND FOCUS**

# HOW DID WE GET THERE



Things happened,  
lets fix them

**REACTIONARY**

8

Reply hazy, try again  
later

**PRECAUTIONARY**



"The best time to start  
was last year"

**TIME AND FOCUS**

# DOCUMENT YOUR ALERTS





# DOCUMENT YOUR ALERTS

Take a holistic  
inventory of what you  
have today


Not just fresh  
memory but dig deep





# DOCUMENT YOUR ALERTS

Look through your  
systems

- messaging
  - monitoring system
  - log management
  - on-call paging
  - codebase
  - etc...
- 





# DOCUMENT YOUR ALERTS

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1	Source	Alert Name	Description	Status	JIRA Link	Last Triggered	Keep?	Recommended Changes	Channel	Slack Link	Source Link	Frequency	Whose Alerted	How Alerted	When Alerted
2	Datadog	High CPU Usage	ECS Task CPU Usage higher than X% over past 15 minutes	Reviewing	<a href="#">Link</a>	1/13/2025 10:00 AM	Yes		alerts	<a href="#">Link</a>	<a href="#">Link</a>	Hourly	All cloud ops	Squadcast	When happens
3	Sentry	X Errors in Past 10 Minutes	Total errors in sentry for project over x	Done	<a href="#">Link</a>	1/13/2025 10:00 AM	Yes		dev	<a href="#">Link</a>	<a href="#">Link</a>	Every weekday	All dev	Slack	When happens

[Sample Link](#)



# DOCUMENT YOUR ALERTS


---

Build out runbooks  
for your alerts





**MEASURE WHAT  
MATTERS**



Just because you  
have the  
measurement doesn't  
mean it's useful

---

**MEASURE WHAT  
MATTERS**






Are your vitals easily  
obtainable

Know your system  
architecture

---

**MEASURE WHAT  
MATTERS**






Some measurements  
you don't know you  
need until after the  
fact

---

**MEASURE WHAT  
MATTERS**





Inventorying our  
alerts allowed us to  
see what overlapped  
and could be reduced

---

**MEASURE WHAT  
MATTERS**





# EVALUATE THE URGENCY






# EVALUATE THE URGENCY

---

Make sure your  
measurements are  
actually right to  
indicate someone  
needs to do  
something






# EVALUATE THE URGENCY

---

Can you configure  
your paging system  
or your monitors to  
work when you need  
them to?






# EVALUATE THE URGENCY

Can you set work and  
non-work hours?

Do you have off  
hours where an alert  
is not necessary?






# EVALUATE THE URGENCY


---

Is the action that failed necessary to alert about or can it be a lower priority notification?





**DON'T LET THE  
TOOLS DICTATE  
YOUR STRATEGY**




Before throwing  
money at a tool,  
figure out what your  
ideal world looks like

---

**DON'T LET THE  
TOOLS DICTATE  
YOUR STRATEGY**








Evaluate your toolset  
Does a higher edition  
contain what you  
need?  
Does a competitor do  
what you need?

---

**DON'T LET THE  
TOOLS DICTATE  
YOUR STRATEGY**





System doesn't have  
an integration with  
your paging solution...  
Build One

---

**DON'T LET THE  
TOOLS DICTATE  
YOUR STRATEGY**





Create custom  
system alerts in your  
application code

---

**DON'T LET THE  
TOOLS DICTATE  
YOUR STRATEGY**



**THINK SMALLER**






# THINK SMALLER

Prove out your  
hypothesis first

Don't just flip a  
switch and enable the  
world






# THINK SMALLER

---


Focus on a the  
highest impact first,  
iterate if necessary  
and move on





**BE FRUGAL BUT  
DON'T BE CHEAP**






Don't have the metric  
you need, figure a  
way to add it even if it  
means a couple  
bucks a month

---

**BE FRUGAL BUT  
DON'T BE CHEAP**






Reallocate money to  
another tool that has  
more value than  
another

---

**BE FRUGAL BUT  
DON'T BE CHEAP**





Don't need top tier  
enterprise solution,  
but maybe that next  
level solves your  
issues

---

**BE FRUGAL BUT  
DON'T BE CHEAP**





**PRACTICE  
MAKES BETTER**

---






# PRACTICE MAKES BETTER

---

Things run smoothly  
and then all of a  
sudden, 3 AM phone  
calls happen






# PRACTICE MAKES BETTER

---

Schedule time to  
complete mock  
incidents






# PRACTICE MAKES BETTER

---

Chaos experiment,  
not just a devops  
check but also a  
good way to test your  
processes



# PRACTICE MAKES BETTER








# PRACTICE MAKES BETTER

---

Build automation to  
not make it on-calls  
responsibility





**ALERT THE  
RIGHT FOLKS**



Associate alerts to  
the right owners, not  
just anyone who is  
on-call



# **ALERT THE RIGHT FOLKS**




Have the proper  
escalations in place

Have everyone who  
can fix something in  
your alerting system

---

**ALERT THE  
RIGHT FOLKS**





One person who  
knows how to fix it  
all? Stop that and  
start sharing the  
wealth. Bus factor is  
a real thing.


---

**ALERT THE  
RIGHT FOLKS**



# WHO KNOWS YOUR SYSTEM BEST?






# WHO KNOWS YOUR SYSTEM BEST?

---

Work with Customer  
Support, Sales,  
Customer Success,  
etc... to understand  
vital flows






# WHO KNOWS YOUR SYSTEM BEST?

---

A process you think  
is important and  
crucial to the system  
might actually not  
matter to the end  
user








# WHO KNOWS YOUR SYSTEM BEST?


---

Understand your user  
usage patterns, both  
how they use the app  
and when they use it






**BUILD INTO  
SDLC**



Make monitoring and  
alerting part of your  
process, not an  
afterthought

**BUILD INTO  
SDLC**





If you're not  
measuring your  
system, how do you  
know it's impact?

**BUILD INTO  
SDLC**



**DON'T JUST FIX  
IT, PREVENT IT**






# **DON'T JUST FIX IT, PREVENT IT**

---

Incidents happen,  
after they happen  
make sure they don't  
bite you again






# DON'T JUST FIX IT, PREVENT IT

Build tests for what  
broke

Take your post  
mortems seriously  
and have traceable  
takeaways






# **DON'T JUST FIX IT, PREVENT IT**

---

Act on your  
takeaways ASAP  
from a post mortem,  
don't let them get  
buried in a backlog





# TOOL SHOUTOUTS



## Checkly

Synthetic monitoring that allows you to do API heartbeat style checks as well as Playwright E2E tests



## Squadcast

All in one incident management tool



## Datadog

Observability service for monitoring servers, databases, and other aspects of your products



## Sentry

Application performance monitoring and error tracking for all of your systems



# THANK YOU

---

Ways to reach me:

clymerrm

