

Putting AI to Work

What the heck is an MCP server?

About Me



- Charlotte, NC
- Developer Relations
- Open Source Nerd
- Not an AI Expert

Our Agenda

- Introductions
- Timeline of AI
- What is MCP?
- Experimenting
- More Resources
- Q & A

Introductions

Tell us About You

- Your name
- What you do
- What was the first album you ever purchased?
(or listened to)

The Current AI Moment

Where did we come from?

1950s - 1970s (AI that Reasons)

the birth of AI, symbolic AI, reasoning

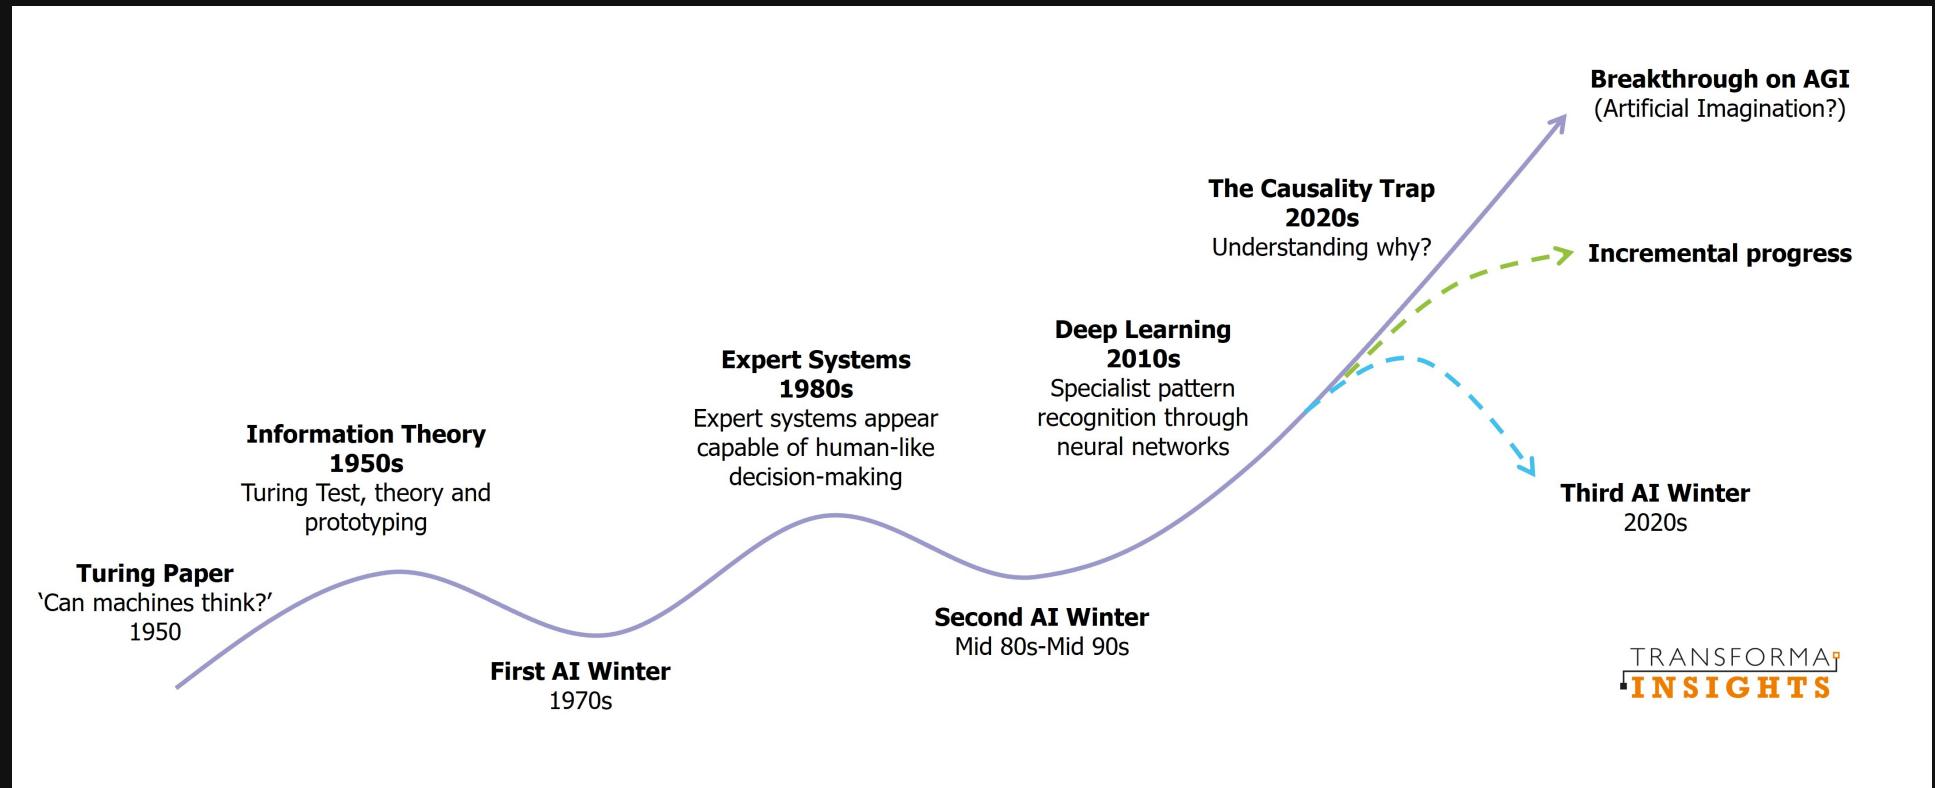
1980s - 1990s (AI Winter)

decision engines, expert systems, niche uses

2000s - 2010s (AI that Predicts)

ML/DL, faster hardware, speech and image recognition





Where are we right now?

2020s (AI that Creates)

LLMs/Generative AI - text, images, code, etc.

2023 - Present (AI that Acts)

Agentic AI - operationalizing LLMs

What is MCP?

How did it come to be?



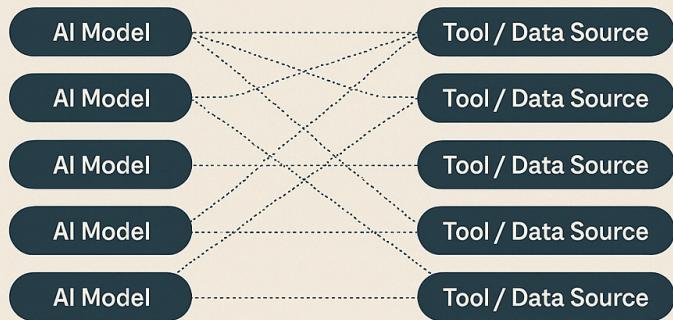
10⁸–10⁹

APIs and Programmatic Interfaces

The Fundamental Problem MCP Solves

NxM Problem

For every AI model that needed to connect with every external tool or data source, a unique integration had to be built and maintained.



1 Development inefficiency

Teams spent countless hours building and maintaining separate integrations

2 Inconsistent implementations

Different developers implemented similar functions in wildly different ways

3 Innovation barriers

Organizations struggled to scale their AI applications

4 Siloed capabilities

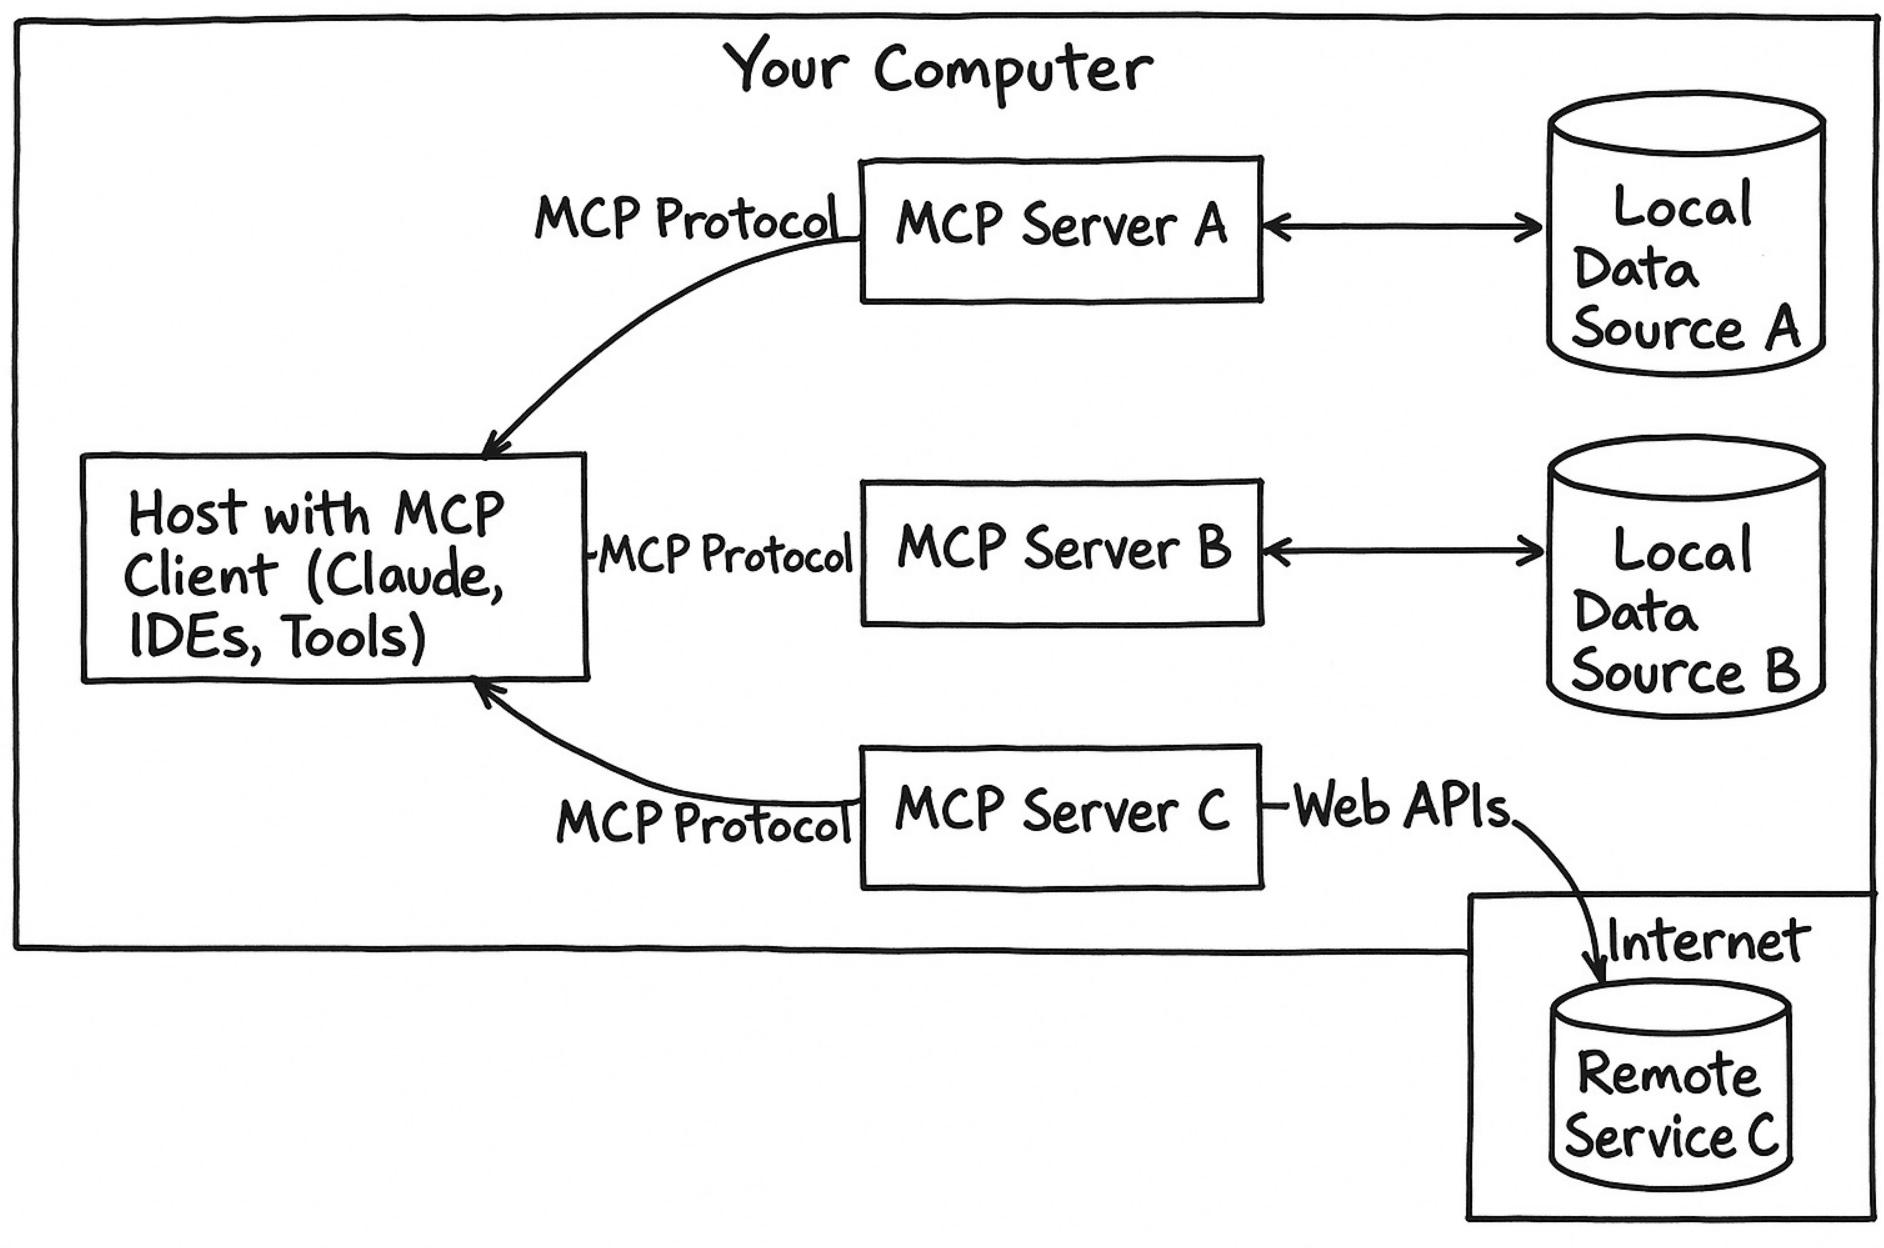
Advanced integration work remained trapped within specific systems

The NxM Problem

- Many competing AI tools and many, many, many external tools
- Developed by Anthropic, announced Nov 2024
- Improved upon ChatGPT's plug-in framework / function-calling API (2023)
- Aimed at addressing the NxM data integration problem

Model-Context-Protocol

- A standard way for LLM tools (Claude, Copilot, ChatGPT) to talk to the outside world.
- Minimizes the need for custom integrations.
- Think of it kind of like ODBC or a USB-C port for AI applications.



Key Pieces of MCP

Hosts + Clients

The AI-powered applications that communicate with MCP Servers using dedicated MCP Clients.

Servers

Expose data and actions to AI Host's clients (often via plugins or adapters)

Protocol

It's JSON-RPC under the hood.

Layers

Transport Layer

Outer layer; manages communication and authentication

Data Layer

Inner layer; defines primitives (tools, resources, prompts, etc.)

Server Primitives

Tools ("What I can do")

Executable functions (API Calls, file operations, DB procedures)

Resources ("What I know")

Data sources (file contents, DB records, API responses)

Prompts (structured interactions)

Reusable templates for structured interactions (plan a vacation, draft an email, etc.)

Client Primitives

Sampling

Lets servers request LLM completions from AI application (ex: Pick the best flight of these)

Elicitation

Lets servers request additional info from users (ex: Ask for the user's seat preference)

Logging

Lets servers send log messages to clients for debugging/monitoring.

Notifications

- Real-time updates
- Tool changes or new functionality
- Clients don't need to poll for changes
- Lets AI apps adapt to changing contexts

Why is MCP a big deal?

- Devs: Reduced development time
- Users: "Smarter" AI for end users
- Apps: Plug-and-play functionality
- Everyone: Cross-platform AI synthesis

Let's Tinker a Bit

More Resources

- <https://modelcontextprotocol.io>
- <https://mcp.so>
- <https://github.com/umbraco/Umbraco-CMS-MCP-Dev>
- <https://www.youtube.com/@IBMTechology>
- Model Context Protocol with Tim Berglund

Final Thoughts

Final Thoughts

Reason → Predict → Create → Act → ???

Final Thoughts

Reason → Predict → Create → Act → ???

-  Collaborate
-  Adapt
-  Invent

Q & A